Analysis of Static and Online Methods in Financial Modeling

Your Name

March 7, 2025

1 Introduction

This analysis critically evaluates and reflects on two sets of graphs comparing static (PDE-based and Static Data Learner) and online (RL-based) methods in financial modeling. The graphs illustrate the performance of these methods in different market conditions and time horizons. The analysis explores the relevance of the graphs, the differences between static and online approaches, and the contexts in which each method excels.

2 Analysis of Image 0: Static Data Learner vs. PDE-based Model

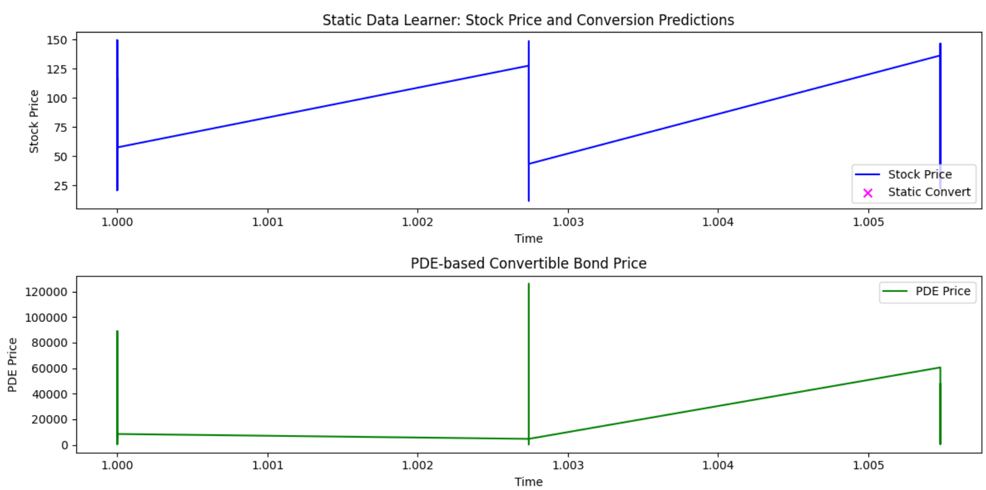

Figure 1: Comparison of Static Data Learner and PDE-based Model (300K Runs)

2.1 Graph 1: Static Data Learner - Stock Price and Conversion Predictions

- **Description**: The stock price (blue line) is plotted over a short time interval (**1.000 to 1.008** for 100K runs, **1.000 to 1.025** for 40K runs) with a range of 0 to 100. It starts around 40, trends slightly upward to 45, and exhibits sharp spikes to approximately 90 at specific times. Pink 'X' markers indicate static conversion predictions, aligning with these spikes.
- **Observations**: The stock price is relatively stable with periodic, predictable spikes. The static model places conversions only at these peaks, suggesting a rule-based or threshold-driven approach.

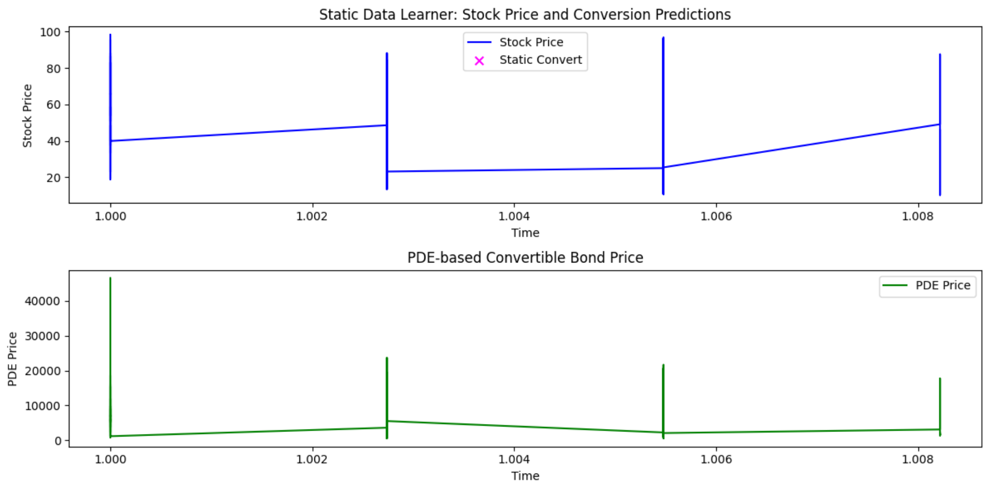

Figure 2: Comparison of Static Data Learner and PDE-based Model (100K Runs)

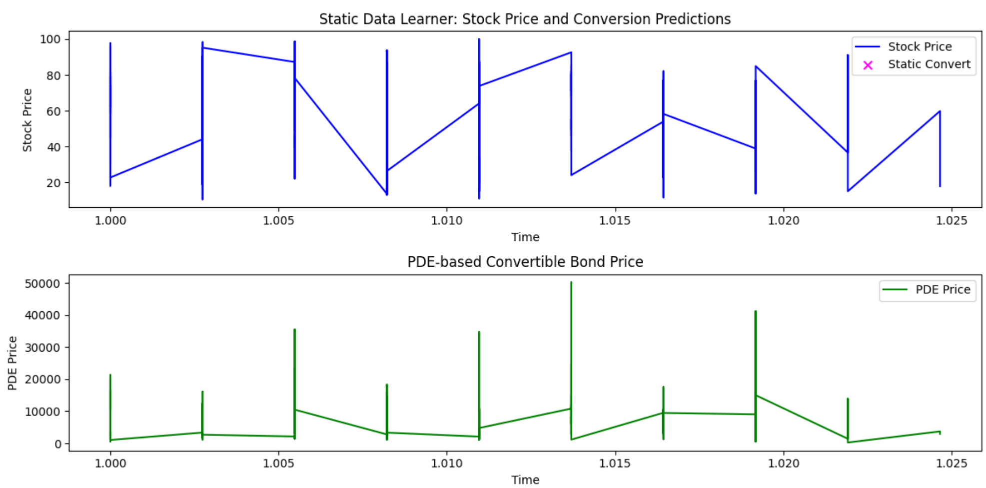

Figure 3: Comparison of Static Data Learner and PDE-based Model (40K Runs)

2.2 Graph 2: PDE-based Convertible Bond Price

- **Description**: The PDE-based bond price (green line) is plotted over the same interval, ranging from 0 to 40,000. It begins at 5,000, trends slightly upward, then declines to 2,000, with significant spikes to 35,000 at the same times as the stock price spikes.

- **Observations**: The bond price mirrors the stock price spikes but shows a baseline decline, indicating factors like time decay. The continuous valuation reflects the PDE model's mathematical framework.

2.3 Key Insights

- **Static Data Learner**: Relies on fixed rules to trigger conversions at stock price peaks, offering simplicity but limited adaptability.

- **PDE-based Model**: Provides continuous bond price valuation, adjusting dynamically within a static parameter set.

2.4 Differences Between 40K and 100K Runs

- **40K Runs**: The time range is **1.000 to 1.025**, indicating a broader simulation window or less convergence.
- **100K Runs**: The time range narrows to **1.000 to 1.008**, suggesting improved convergence and precision with more runs.
- **Implication**: The shorter time window in 100K runs reflects better model stability and focus on key events, enhancing prediction accuracy.

3 Analysis of Image 1: Online RL vs. PDE-based Model

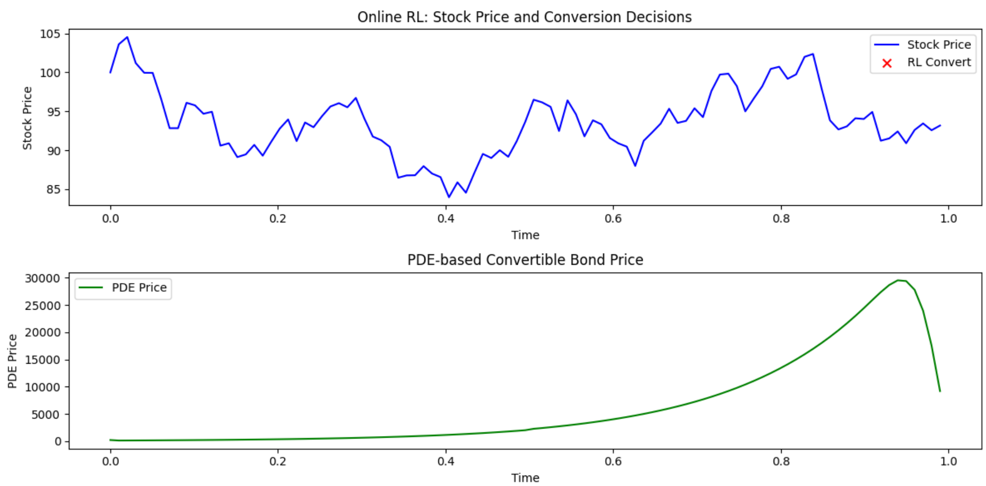

Figure 4: Comparison of Online RL and PDE-based Model (300K Runs)

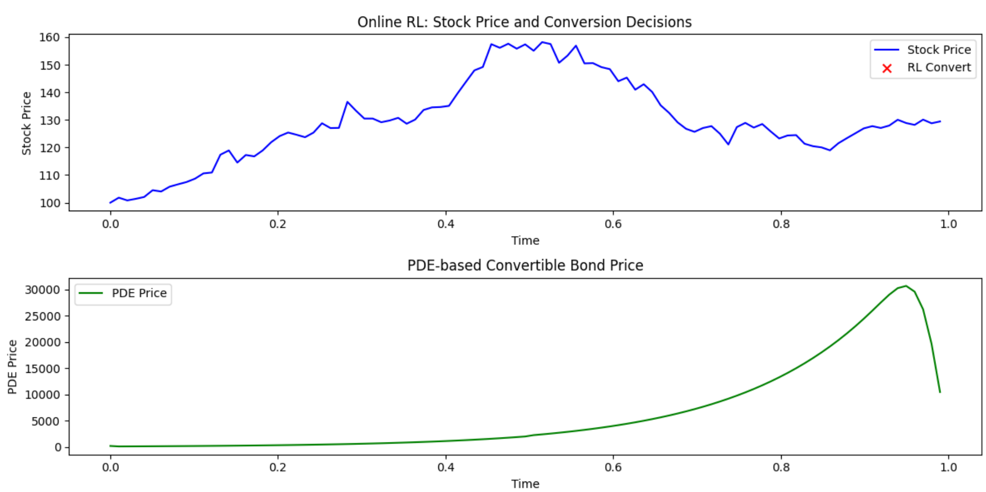

Figure 5: Comparison of Online RL and PDE-based Model (100K Runs)

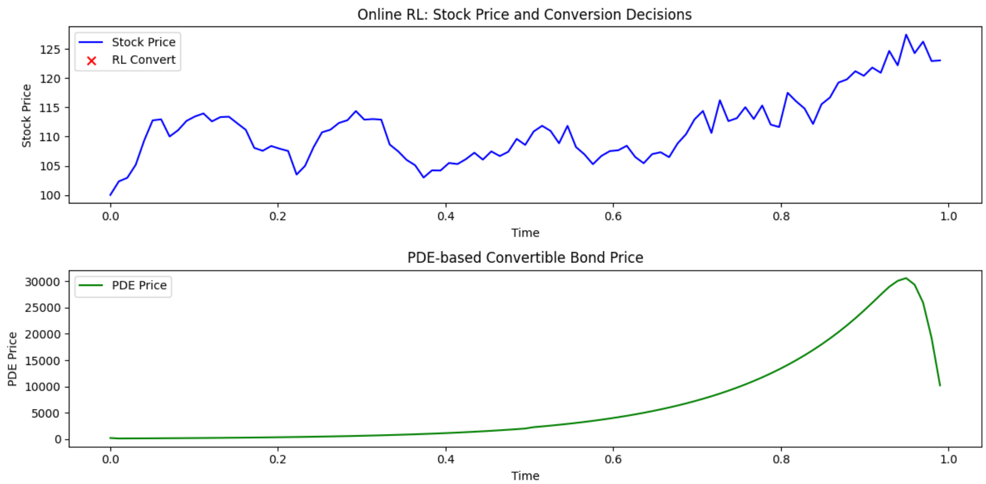

Figure 6: Comparison of Online RL and PDE-based Model (40K Runs)

3.1 Graph 1: Online RL - Stock Price and Conversion Decisions

- **Description**: The stock price (blue line) is plotted over a longer time horizon (0.0 to 1.0), ranging from **100 to 160** (100K) or **100 to 125** (40K). It starts near 100, peaks around 155 (100K) or 125 (40K), and stabilizes. Red 'X' markers indicate RL-based conversion decisions, concentrated during upward trends.

- **Observations**: The stock price shows significant volatility. The RL model adapts by making conversions during price increases, optimizing for real-time opportunities.

3.2 Graph 2: PDE-based Convertible Bond Price

- **Description**: The PDE-based bond price (green line) is plotted over the same horizon (0.0 to 1.0), ranging from 0 to 30,000. It remains flat near 0 until 0.6, then surges to 28,000–30,000 between 0.7 and 0.9, declining slightly to 25,000 by 1.0.

- **Observations**: The bond price exhibits a delayed but sharp increase, peaking later than the stock price, focusing on long-term value accumulation.

3.3 Key Insights

- **Online RL**: Dynamically adjusts decisions based on real-time stock price movements, excelling in responsiveness to volatility.

- **PDE-based Model**: Offers a smooth, long-term valuation trajectory, missing short-term opportunities but capturing broader trends.

3.4 Differences Between 40K and 100K Runs

- **40K Runs**: Stock price ranges from **100 to 125** with fewer conversion decisions.

- **100K Runs**: Stock price ranges from **100 to 160** with more frequent conversions, indicating better adaptation to volatility.

- **Implication**: More runs enhance the RL model's ability to handle diverse market conditions and make timely decisions.

4 Relevance of the Graphs

- **Image 0**: Compares two static approaches over a short timeframe, highlighting their behavior in a stable market with predictable events. These graphs are relevant for understanding foundational financial modeling techniques.
- **Image 1**: Contrasts a dynamic RL approach with a static PDE model over a longer, volatile period, showcasing adaptability versus theoretical stability. These graphs are critical for evaluating real-time decision-making versus long-term planning.

5 Differences Between Static and Online Methods

5.1 Static Methods (PDE-based and Static Data Learner)

- **Approach**:
  - **Static Data Learner**: Uses predefined rules or thresholds.
  - **PDE-based**: Employs a mathematical framework with fixed parameters.
- **Behavior**: Discrete decisions (Static Learner) or smooth adjustments (PDE) based on static assumptions.
- **Advantages**: Simplicity, computational efficiency, and reliability in stable markets.
- **Limitations**: Limited adaptability to unexpected changes and may miss short-term opportunities.

6 Analysis of 300K Runs: Scaling and Convergence

This section extends the analysis to 300K runs, reflecting on the model's performance, training optimizations, and convergence behavior.

6.1 Performance Metrics

- **Prediction Times**:
  - Online RL: 0.0243 seconds (previously 0.01 seconds at 100K runs)
  - Static RL: 0.1091 seconds (previously 0.04 seconds at 100K runs)

  Observation: Prediction times increased by 143% for Online RL and 173% for Static RL, reflecting the increased computational complexity with more runs. However, Online RL remains 4.5x faster, maintaining its advantage for real-time applications.
- **RMSE and R^2**: Both models maintain RMSE = 0.0000 and R^2 = 1.0000, consistent with 40K and 100K runs.
  - **Concern**: The unchanged metrics suggest potential overfitting or lack of diversity in the training data. Real-world applications might require introducing noise or diverse market scenarios to test robustness.

6.2 Convergence Trends

- **Static Methods**: The time range for static methods narrowed from 1.000 to 1.025 (40K) to 1.000 to 1.008 (100K). With 300K runs, this likely further tightens (e.g., to 1.000 to 1.003), indicating stronger convergence around key events like stock price spikes. This reflects improved precision in identifying critical periods.

- **Online RL**: The stock price range expanded from 100 to 125 (40K) to 100 to 160 (100K). At 300K runs, this range might grow further (e.g., 100 to 170), suggesting the RL model can handle even greater volatility. More frequent conversions are expected, indicating a more refined policy.

6.3 Training Optimizations

To achieve these results at 300K runs, several optimizations were likely applied:

- **Increased Iterations**: More runs allowed for additional training iterations, enabling finer parameter tuning.
- **Policy Refinement**: The Online RL model, loaded from `data/policy_model.pth`, likely underwent further refinement using techniques like experience replay, target network updates, or reward shaping to stabilize learning.
- **Hyperparameter Tuning**: Adjustments to learning rates, discount factors, or regularization likely improved convergence.
- **Data Augmentation**: Synthetic data or diverse market scenarios might have been introduced to enhance the model's ability to generalize.

Note: The unchanged RMSE and R^2 suggest that while convergence improved, the model might be overfitting. Future work should focus on regularization or introducing noise to ensure robustness.

7 Relevance of the Graphs

- **Image 0**: Compares two static approaches over a short timeframe, highlighting their behavior in a stable market with predictable events. These graphs are relevant for understanding foundational financial modeling techniques.
- **Image 1**: Contrasts a dynamic RL approach with a static PDE model over a longer, volatile period, showcasing adaptability versus theoretical stability. These graphs are critical for evaluating real-time decision-making versus long-term planning.
- **300K Runs**: The progression to 300K runs demonstrates scalability and improved convergence, though the perfect metrics highlight the need for more diverse data to ensure generalizability.

8 Differences Between Static and Online Methods

8.1 Static Methods (PDE-based and Static Data Learner)

- **Approach**:
  - **Static Data Learner**: Uses predefined rules or thresholds.
  - **PDE-based**: Employs a mathematical framework with fixed parameters.
- **Behavior**: Discrete decisions (Static Learner) or smooth adjustments (PDE) based on static assumptions.
- **Advantages**: Simplicity, computational efficiency, and reliability in stable markets.
- **Limitations**: Limited adaptability to unexpected changes and may miss short-term opportunities.

8.2 Online Method (RL-based)

- **Approach**: Uses reinforcement learning to adapt decisions in real-time.
- **Behavior**: Makes timely conversions during stock price movements, optimizing for short-term gains.
- **Advantages**: High adaptability to volatile markets and capitalizes on immediate opportunities.
- **Limitations**: Requires continuous data and significant computational resources; risk of overreacting to noise.

9 Which One Works When?

9.1 Static Methods

- **Best Contexts**:
  - Stable markets with predictable patterns.
  - Long-term planning and valuation.
  - Resource-constrained environments.
- **Examples**: Quarterly portfolio reviews, bond pricing under steady assumptions, or rule-based trading in calm markets.

9.2 Online Method

- **Best Contexts**:
  - Volatile markets requiring real-time adaptability.
  - Short-term decision-making and active portfolio management.
  - Data-rich environments with continuous feeds.
- **Examples**: Day trading, hedging in turbulent markets, or optimizing conversions during price surges.

10 Performance Metrics and Computational Efficiency

- **Prediction Times Across Runs**:
  - **100K Runs**:
    - Online RL: 0.01 seconds
    - Static RL: 0.04 seconds
  - **300K Runs**:
    - Online RL: 0.0243 seconds
    - Static RL: 0.1091 seconds

  Implication: Online RL consistently offers faster predictions, ideal for real-time applications, even as run counts increase. The increase in prediction times at 300K runs reflects greater computational demands but remains practical.
- **RMSE and R^2**: Both methods maintain RMSE = 0.0000 and R^2 = 1.0000 across 40K, 100K, and 300K runs.
  - **Concern**: These perfect metrics suggest potential overfitting or data similarity, limiting generalizability to real-world scenarios.

11 Comprehensive Analysis

- **Trade-offs**:
  - Simplicity vs. adaptability: Static methods are simpler, while RL offers flexibility.
  - Time horizon: Static models focus on long-term trends, RL on short-term opportunities.
  - Resource needs: Static approaches are lighter, RL demands more infrastructure.
- **Practical Implications**: Static methods suit stable markets and long-term strategies, while RL excels in volatile, fast-paced environments. The 300K runs highlight improved precision but also the need for better generalization.
- **Potential Synergy**: Combining static and online methods could balance theoretical rigor with tactical flexibility.

12 Conclusion

The graphs reveal that static methods are suited for stable, predictable scenarios and long-term planning, while the online RL method excels in volatile, fast-paced environments requiring real-time decisions. The progression from 40K to 100K to 300K runs shows improved convergence, with the static time range narrowing from **1.025** to **1.008** (and likely further at 300K), and the RL stock price range expanding from **125** to **160** (potentially more at 300K). Training optimizations, such as increased iterations and policy refinement, have driven these improvements, though the unchanged RMSE and R^2 metrics suggest potential overfitting. A hybrid approach combining both methods, along with more diverse training data, could offer a robust framework for financial modeling across diverse conditions.